

DIALECTIC: A Multi-Agent System for Startup Evaluation

Anonymous ACL submission

Abstract

Venture capital (VC) investors face a large number of investment opportunities but only invest in few of these, with even fewer ending up successful. Early-stage screening of opportunities is often limited by investor bandwidth, demanding tradeoffs between evaluation diligence and number of opportunities assessed. To ease this tradeoff, we introduce DIALECTIC, an LLM-based multi-agent system for startup evaluation. DIALECTIC first gathers factual knowledge about a startup and organizes these facts into a hierarchical question tree. It then synthesizes the facts into natural-language arguments for and against an investment and iteratively critiques and refines these arguments through a simulated debate, which surfaces only the most convincing arguments. Our system also produces numeric decision scores that allow investors to rank and thus efficiently prioritize opportunities. We evaluate DIALECTIC through backtesting on real investment opportunities aggregated from five VC funds, showing that DIALECTIC matches precision of real VCs in predicting startup success.

1 Introduction

The global venture capital (VC) industry is expanding rapidly alongside intensified competition for attractive deals. The market is projected to grow from USD 337 billion in 2024 to USD 1.46 trillion by 2033, a compound annual growth rate of 17.6% ([IMARC Group, 2024](#)). Entrepreneurial activity has also surged, with annual U.S. business formations rising from 3.5 million in 2019 to 5.2 million in 2024, an increase of nearly 40% ([U.S. Census Bureau, 2025](#)). Traditional VC decision-making processes are challenged in this setting. Investors face high time pressure and information overload, both associated with suboptimal decision ([Zacharakis and Shepherd, 2001](#)). These conditions have increased interest in computational approaches for scalable investment evaluation.

Among these approaches, machine learning methods have emerged as a promising direction. Prior studies demonstrate strong predictive performance and, in some cases, even surpass human investors ([Antretter et al., 2019; Retterath, 2020; Zacharakis and Shepherd, 2001; Arroyo et al., 2019; Dellermann et al., 2021; Sharchilev et al., 2018](#)). Yet, these non-iterative models diverge from how investment decisions are formed by human VCs. In practice, conviction emerges through iterative hypothesis formation, challenge, and refinement as new information appears ([Chong and Tuckett, 2014](#)).

Recent advances in *large language model* (LLM) orchestration tools enable iterative and interpretable reasoning. Frameworks such as *LangChain* ([Chase, 2022](#)) support the decomposition of complex tasks, the generation of intermediate conclusions, and the iterative refinement of responses while making the underlying logic explicit. They allow multi-step reasoning and dialectical interaction, a setup in which LLMs can articulate arguments, generate counterpoints, and produce transparent reasoning traces.

This paper introduces *Decision Iteration with Argument-Level Evidence and Counter-Thinking for Investment Conclusions* (**DIALECTIC**), an LLM-based system that models iterative and argumentative elements of venture evaluation. Our system draws on principles of dialectical reasoning, an approach shown to be effective for complex, unstructured problems that benefit from structured confrontation of differing perspectives ([Jarpaphirun and Zahedi, 2007](#)). The contributions of this work are:

- A structured LLM reasoning system that models how investors build and refine investment theses through argumentation.
- An empirical evaluation demonstrating predictive performance in venture screening.

083 Overall, the proposed system brings data-driven
084 VC methods closer to industry practice. Furthermore,
085 it enables the process of iterative argumentation
086 in early-stage screening, which has traditionally
087 been restricted to later stages of the funnel
088 due to limited investor bandwidth. This shift al-
089 lows investors to apply iterative reasoning earlier
090 in the process, improving both diligence quality
091 and screening efficiency.

092 2 Related Work

093 Prior studies propose different machine learning
094 approaches to predict startup outcomes, drawing
095 on public data sources such as *Crunchbase* (Arroyo
096 et al., 2019; Źbikowski and Antosiuk, 2021; Ret-
097 terath, 2020), *Twitter* (Antretter et al., 2019), web
098 data (Sharchilev et al., 2018), and *Google Search*
099 (Gavrilenco et al., 2023), and often reporting
100 promising prediction accuracy (see Table 4 in the
101 Appendix for an overview). Most studies trained
102 gradient tree boosting models (e.g., *XGBoost*)
103 (Corea et al., 2021; Arroyo et al., 2019; Źbikowski
104 and Antosiuk, 2021; Retterath, 2020) and inter-
105 preted predictions using feature-importance rank-
106 ings with features such as geography, industry,
107 or founder background (Źbikowski and Antosiuk,
108 2021; Sharchilev et al., 2018; Gavrilenco et al.,
109 2023).

110 Some newer studies have used LLMs to extract
111 structured features or embeddings from unstruc-
112 tured data, while still resorting to machine learning
113 models such as *XGBoost* for prediction (Ozince
114 and Ihlamer, 2024; Maarouf et al., 2025). Xiong
115 and Ihlamer (2023) used LLMs to assess founder-
116 idea fit, also providing pro and contra arguments
117 for interpretability. In follow-up work (Xiong et al.,
118 2024), they focus on extracting traits associated
119 with successful entrepreneurs. Both studies look at
120 individual founders rather than startup companies.

121 Beyond VC, LLM-based decision-making frame-
122 works have been proposed for fields such as busi-
123 ness or finance. *DeLLMa* combines LLMs with
124 decision-theoretic reasoning (Liu et al., 2025),
125 while *STRUX* extracts facts from companies’ earn-
126 ings calls and produces weighted pro and contra
127 aspects for buy or sell decisions (Lu et al., 2025).

128 A promising approach to improving LLM rea-
129 soning is the introduction of multi-agent systems
130 (Han et al., 2024). Instead of relying on a sin-
131 gle model, several LLMs interact through collab-
132 oration, debate, or specialization. In adversarial

133 or collaborative debating, agents defend opposing
134 stances and a separate judge model or heuristic eval-
135 uates the quality of their arguments (Chan et al.,
136 2023; Liang et al., 2024).

137 3 DIALECTIC

138 DIALECTIC is inspired by how real VCs make in-
139 vestment decisions. They collect information about
140 a startup, form narrative investment hypotheses,
141 and refine these hypotheses through debate with
142 other VCs until making a decision. DIALECTIC
143 proceeds in three phases: **fact collection**, **reason-
144 ing**, and **decision-making**. During fact collection,
145 DIALECTIC gathers factual knowledge about a
146 company and organizes these facts hierarchically
147 in a question tree. In the reasoning phase, it syn-
148 thetizes raw facts into arguments pro and contra an
149 investment, which it iteratively self-critiques, eval-
150 uates, and refines, letting only the best arguments
151 survive. Finally, it makes a decision based on a
152 comparison of the best pro and contra arguments;
153 see Figure 1 for an illustration.

154 In the following, we formally introduce DI-
155 ALECTIC. Let $X = \{x_i\}_{i=1}^N$ be the set of inves-
156 tigable companies, each described by multiple fea-
157 tures $x_i^{(d)}$, $d = 1, \dots, D$. The goal is to predict the
158 ground truth label $y_i \in \{\text{successful, unsuccessful}\}$
159 signaling whether the company will be successful
160 and should be invested in or not.

161 3.1 Fact Collection Phase

162 We denote the universe of natural-language ques-
163 tions as \mathcal{Q} , the universe of natural-language an-
164 swers as \mathcal{A} , and the set of industries as I . For a
165 given company x in industry $x^{(0)} \in I$, we start by
166 providing DIALECTIC with a set of **seed ques-
167 tions** $Q_0 \subset \mathcal{Q}$. Specifically, we ask four questions
168 about the general company, team, product, and mar-
169 ket (see Appendix B.1 for details) to cover the main
170 aspects typically considered by VC investors (Ret-
171 terath, 2020). Inspired by *ProbTree* (Cao et al.,
172 2023) and *Socratic Questioning* (Qi et al., 2023),
173 we define two LLM-based agent operations:

- 174 • The **decomposer** $Q : Q_0 \times I \rightarrow \mathcal{Q}^*$
175 takes a seed question $q \in Q_0$ and hierar-
176 chically decomposes it into a finite set of
177 M_q sub-questions relevant in the industry,
178 thus creating an industry-specific **question
179 tree**¹ $\{q_l\}_{l=1}^{M_q} = Q(q, x^{(0)})$ with the decision-
180 relevant questions that should be answered.

¹For simplicity, we do not explicitly model the hierarchical

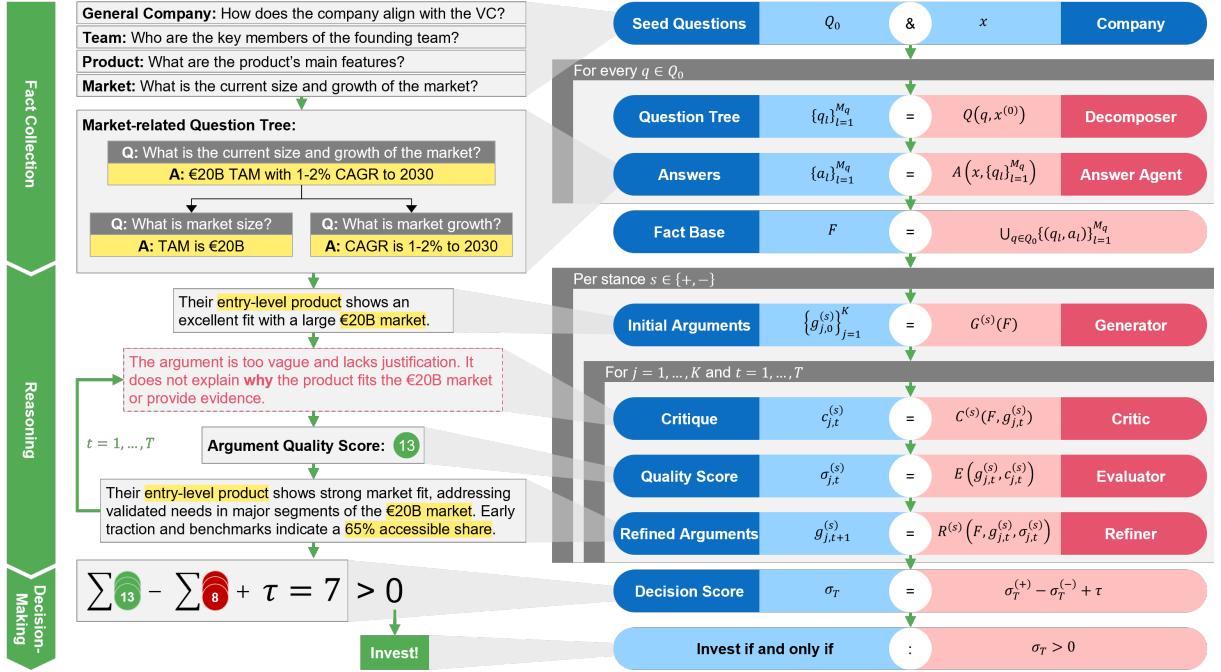


Figure 1: Overview of the DIALECTIC method. The right side shows the flow of operations. Agents are shown in red, agent inputs/outputs are shown in blue, and loops are shown in green. The left side illustrates the key outputs of the agents.

- The **answer agent** $A : X \times \mathcal{Q}^* \rightarrow \mathcal{A}^*$ looks at the company features $x^{(d)}$ and uses them to generate answers $\{a_l\}_{l=1}^{M_q}$ to all questions in the tree. It also has access to a web search tool that it can use agentically. Like *ProbTree* (Cao et al., 2023), it answers question trees in a post-order traversal, aggregating answers from child nodes when generating answers for parent nodes. We provide further details in Appendix B.2.

When executed for all seed questions, the agents produce a rich hierarchically structured **fact base** $F \subset \mathcal{F}$ (\mathcal{F} is the universe of all possible question-answer pairs) about the company x :

$$F = \bigcup_{q \in Q_0} \{(q_l, a_l)\}_{l=1}^{M_q}$$

3.2 Reasoning Phase

In the reasoning phase, DIALECTIC combines facts (possibly from different question trees) into arguments taking a stance on whether the VC should invest in a company or not. Let $s \in \{+, -\}$ denote the pro or contra stance, \mathcal{G} the universe of natural-language arguments, and \mathcal{C} the universe of

organization of questions in our notation but represent question trees as simple sets. However, our code implementation preserves the full hierarchy.

natural-language critiques of these arguments. We define four LLM-based agent operations:

- The **generator** $G^{(s)} : \mathcal{F} \rightarrow \mathcal{G}^K$ takes the fact base and generates K **arguments** $\{g_j^{(s)}\}_{j=1}^K = G^{(s)}(F)$ per stance s , citing various facts from the fact base. This is inspired by Park et al. (2023) who recursively synthesize observations into higher-level reflections.
- The **critic** $C^{(s)} : \mathcal{F} \times \mathcal{G} \rightarrow \mathcal{C}$ criticizes an argument, producing a **critique** $c_j^{(s)} = C^{(s)}(F, g_j^{(s)})$ against it, possibly also citing facts from the fact base. The critic thereby acts as a *devil’s advocate* (Kim et al., 2024) sparking a debate about the argument.
- The **evaluator** $E : \mathcal{G} \times \mathcal{C} \rightarrow \mathbb{N}$ takes an argument and corresponding critique and judges the convincingness of the argument with a **quality score** $\sigma_j^{(s)} = E(g_j^{(s)}, c_j^{(s)}) \in \mathbb{N}$. Internally, it uses a 14-criteria evaluation scheme based on the argument quality taxonomy by Wachsmuth et al. (2017). See Appendix B.3 for further details.
- The **refiner** $R^{(s)} : \mathcal{F} \times \mathcal{G} \times \mathbb{N} \rightarrow \mathcal{G}$ refines a given argument trying to improve its quality. It produces a **refined argument** $\tilde{g}_j^{(s)} = R^{(s)}(F, g_j^{(s)}, \sigma_j^{(s)})$.

228 As the refinement can be repeated, we use the notation
 229 $g_{j,t+1}^{(s)} = R^{(s)}(F, g_{j,t}^{(s)}, \sigma_{j,t}^{(s)})$ instead, where
 230 the index $t = 1, \dots, T$ denotes the iteration.

231 Starting with an initial set of K_0 arguments
 232 $\{g_{j,0}^{(s)}\}_{j=1}^{K_0} = G^{(s)}(F)$, DIALECTIC iteratively criti-
 233 ques, evaluates, and refines the arguments. It
 234 hereby follows a *survival-of-the-fittest* logic, keep-
 235 ing only the best K_t arguments (the **survivors** S_t)
 236 after each iteration t :

$$S_{t+1}^{(s)} = TopK(\{g_{j,t+1}^{(s)} : g_{j,t}^{(s)} \in S_t^{(s)}\}, K_{t+1}),$$

237 where $TopK(\{\cdot\}, K_{t+1})$ denotes the K_{t+1} ar-
 238 guments with the highest quality scores $\sigma_{j,t+1}^{(s)}$ in
 239 $\{\cdot\}$. With arguments iteratively improving and K_t
 240 decreasing over the iterations, DIALECTIC con-
 241 verges to a narrow selection $S_T = S_T^{(+)} \cup S_T^{(-)}$ of
 242 high-quality pro and contra arguments. This mim-
 243 ics a debate in a VC investment committee where
 244 different members have different stances on the in-
 245 vestment and continue to bring forward arguments
 246 until the room converges to a dominant narrative.

247 3.3 Decision-Making Phase

248 After T iterations of debate, a few dominant argu-
 249 ments for either stance have emerged. To determine
 250 which stance has the better arguments, we look at
 251 the sum of the argument quality scores for all sur-
 252 viving arguments and compare the pro and contra
 253 stances, calculating the **decision score** σ_T :

$$\sigma_T = \sigma_T^{(+)} - \sigma_T^{(-)} + \tau,$$

254 where $\sigma_T^{(s)} = \sum \sigma_{j,T}^{(s)}$ is the sum of the quality
 255 scores of all surviving arguments $g_{j,T}^{(s)} \in S_T^{(s)}$ and τ
 256 is a **decision threshold** capturing VC's preference
 257 for a margin of safety. Finally, DIALECTIC will
 258 decide to invest if and only if $\sigma_T > 0$.

259 3.4 Hyperparameters & Implementation

260 The above definition of DIALECTIC presents three
 261 main hyperparameters: The number of arguments
 262 kept per iteration K_t , the number of iterations T ,
 263 and the decision threshold τ . In our implemen-
 264 tation we set $K_t = 5$ for $t \neq T$ and test different
 265 values of K_T , T , and τ . For the LLM, we use Open-
 266 AI's gpt-5-mini-2025-08-07 (OpenAI, 2025). We
 267 set the temperature parameter to 0.0 for the
 268 answer agent and to 0.5 for all other agents. We
 269 report all used prompts in Appendix C.

4 Evaluation Setup

We evaluate our method in a backtesting experiment by predicting startup success from historic data and benchmarking against real VC investors. Our dataset includes 259 startups that were added to real VCs' watchlists² between January 1, 2021 and December 31, 2021. The VCs considered joining the initial funding rounds (*seed* or *pre-seed*) of these startups, which were raised some time between January 1, 2021 and February 28, 2023.

Dependent variable Similar to prior work (Sharchilev et al., 2018; Gavrilenco et al., 2023), we define a startup as *successful*, if it has subsequently raised a *series A* or later round by September 1, 2025, otherwise as *unsuccessful*. With startup success as the dependent variable, our setup is a binary classification. Among all 259 startups, 25% were successful.

Independent variables To predict startup success, the following features are known for each startup: company name, short and long description, industry domain, team description, website content, and web search results (Table 2 in the Appendix presents descriptions of all features). These features are extracted from the VCs' watchlists, Crunchbase.com, startup homepages, and the Perplexity Sonar API. To prevent *look-ahead bias* (Żbikowski and Antosiuk, 2021), we use historic data snapshots and time filters to ensure all features have been available to the VC at the time of the investment decision. See Appendix A for a detailed description of the dataset and its creation.

Baselines We compare the classification performance of our method against the performance of the real VCs. The VCs invested in 6 of the 259 startups, and 2 of these were successful. Also, we compare against simple input-output (IO) prompting (see Listing 10 in the Appendix for the prompt).

Dataset split We split the data into a validation set with 129 startups and a test set with 130 startups using random stratified sampling, so that both sets have an equal ratio of successful startups and each set includes 3 startups that the VCs invested in.

Metrics To measure the performance of DIALECTIC and its baselines, we primarily look at

²The watchlists comprises data of five different VC funds and the real VCs' performance reported in this paper represents a weighted average across these funds.

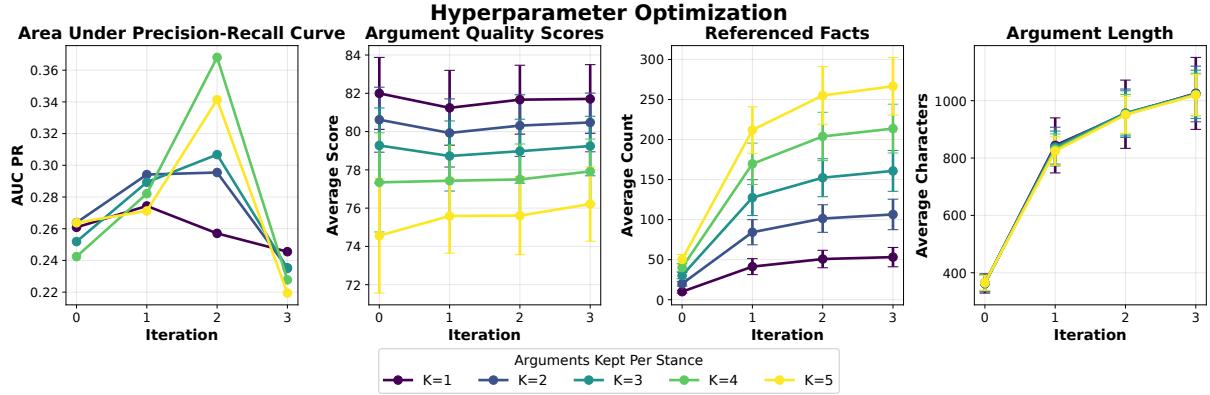


Figure 2: Results from the hyperparameter optimization, showing AUC-PR, raw argument scores, QA pair count, and argument length for different numbers of arguments K_t and iterations T .

precision and recall for different values of the decision threshold τ , as well as the area under curve (AUC) of the precision-recall (PR) line. We also assess the argument quality scores, number of cited facts, length of the arguments, as well as the distribution of decision scores. We first identify a well-performing combination of our hyperparameters T and K_T on the validation set. We then evaluate this configuration on the test set.

5 Results

Our results cover hyperparameter optimization, comparative predictive performance, and an analysis of which facts DIALECTIC uses.

5.1 Hyperparameter Optimization

We optimize the system by varying the number of surviving arguments per side (K_T) and the number of refinement iterations (T). These parameters control how broadly the system explores arguments and how deeply it refines them. Figure 2 summarizes the effect of varying these hyperparameters. Precision–recall performance shows a clear pattern: AUC-PR increases consistently from $T = 0$ to $T = 2$ and declines for $T \geq 3$. The best result occurs at $T = 2$ with $K_2 = 4$, which we use for subsequent experiments.

The diagnostic plots display coherent trends across other measures. Argument quality scores increase with more iterations, with only mild variation across K_t . Referenced facts rise with both parameters, suggesting stronger arguments rely on broader evidence. Argument length jumps from $T = 0$ to $T = 1$, driven by the introduction of structured justification, and grows more slowly thereafter as later iterations add elaboration rather than

new insights. Because length and referenced facts increase smoothly while AUC-PR declines only at $T \geq 3$, the performance drop likely reflects over-refinement effects (e.g., redundancy or drift) rather than simple argument inflation.

5.2 Predictive Performance Against Baselines

Precision and Recall of DIALECTIC vs. Baselines

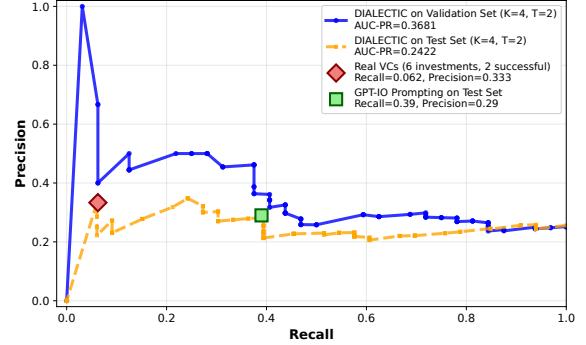


Figure 3: Precision and recall of DIALECTIC across all possible values of the decision threshold τ in comparison to the human VCs and GPT IO prompting baseline

Figure 3 reports the predictive performance of DIALECTIC on the validation set and the held-out test set. The system attains an AUC-PR of 0.2422 on the test set, with precision comparable to human investors and the GPT-IO prompting baseline. Performance is higher on the validation set, where it achieved higher AUC-PR and even outperformed the real VCs. The operating points in Figure 3 show that, in the high-precision, low-recall region, the system behaves similarly to the baselines. Unlike the baselines, it produces a full ranked frontier rather than a single operating point, which allows practitioners to choose a decision threshold tailored to screening capacity. Overall, DIALECTIC is

369 comparable to baselines in predictive performance
 370 while offering a full decision frontier and ranking
 371 rather than a single operating point.

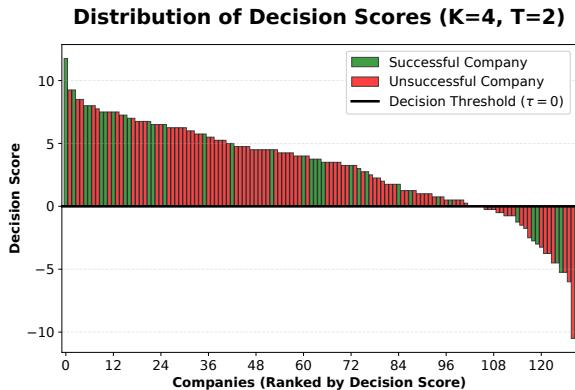


Figure 4: Distribution of decision scores.

372 Figure 4 shows the distribution of decision
 373 scores for the best-performing configuration.
 374 Successful companies cluster at the top of the ranking,
 375 (left end of the plot), while unsuccessful ones ap-
 376 pear toward the bottom (right side of the plot). This
 377 separation shows that higher scores correspond to
 378 a higher likelihood of success. In practice, select-
 379 ing a threshold simply involves choosing a cutoff
 380 along this ranking. A higher threshold prioritizes
 381 the strongest opportunities and filters out most low-
 382 scoring cases.

383 5.3 Evidence Utilization

Aspect	Usage	Availability	Ratio
General	34.40%	34.88%	0.986
Team	20.52%	20.17%	1.017
Product	29.77%	29.43%	1.011
Market	15.31%	15.51%	0.987

384 Table 1: Utilization of factual evidence in arguments.
 385 Aspect refers to the question trees built for the four seed
 386 questions. Usage measures the share of all cited facts,
 387 availability measures the relative size of question trees,
 388 and ratio is the ratio of usage and availability.

389 Table 1 summarizes how the model uses differ-
 390 ent evidence categories when generating factual
 391 references. General company and product infor-
 392 mation dominate, accounting for nearly 65% of all
 393 references, mirroring their share in the fact base.
 394 The “ratio” column reports the ratio between how
 395 often an aspect is referenced and its relative rep-
 396 resentation in the fact base. Values slightly above one

397 (team: 1.017; product: 1.011) indicate that these
 398 aspects are referenced more frequently than their
 399 availability alone would predict. Market informa-
 400 tion is slightly under-used (efficiency < 1). Overall,
 401 ratios cluster near one, indicating proportional use
 402 of available evidence. Notably, team-related evi-
 403 dence receives the highest relative usage, which
 404 aligns with established findings that investors fre-
 405 quently prioritize founder and team attributes when
 406 forming investment judgments (Gompers et al.,
 407 2020).

408 6 Conclusion

409 This paper introduced DIALECTIC, an LLM-based
 410 multi-agent system for early-stage startup screen-
 411 ing. The system integrates fact extraction, argu-
 412 ment generation, iterative critique, and scoring into
 413 one pipeline. Evaluated on an industry-sourced
 414 dataset, the system achieves performance compa-
 415 rable to human investors while producing inter-
 416 pretable argument structures.

417 A central contribution of the approach is the
 418 introduction of iterative argumentation at the be-
 419 ginning of the investment funnel. Since investor
 420 bandwidth limits such deliberation during initial
 421 screening, it usually occurs later in the funnel. En-
 422 abling it earlier provides a structured foundation
 423 for preliminary assessments and supports reasoning
 424 under uncertainty.

425 Operationally, the system reduces time to initial
 426 assessment and produces a ranking when deal vol-
 427 ume exceeds human screening capacity. It supports
 428 both fixed-threshold (returning only companies ex-
 429 ceeding a certain decision score) and fixed-quantity
 430 (returning only the top N companies according to
 431 their decision scores) screening modes, reflecting
 432 constraints encountered in practice. Since miss-
 433 ing strong opportunities is costlier than evaluat-
 434 ing weak ones, the ranking mechanism aligns with
 435 recall-oriented objectives common in top-of-funnel
 436 screening. The generated arguments and evidence
 437 also support later stages such as due diligence or
 438 memo preparation.

439 7 Limitations

440 Venture capital is a domain of high uncertainty and
 441 low signal-to-noise ratios. This is also reflected by
 442 the relative instability of our results. Hyperparam-
 443 eter changes and sample composition can have a
 444 considerable impact on model performance, high-
 445 lighted by the difference in AUC-PR seen between

441 the validation and the test sets.

442 In investment decision tasks, avoiding look-
443 ahead bias is a critical priority. We had to sac-
444 rifice more than 90% of the original companies in
445 our dataset as a result of our efforts to minimize
446 look-ahead bias and ensure feature completeness,
447 leaving only 259 companies in the final dataset (see
448 Appendix A). Similarly, this also restricted the cor-
449 responding baseline of real VCs' performance to
450 data of only 6 invested companies.

451 While we took extensive measures to minimize
452 look-ahead bias, a minor risk of bias remains. As
453 we had access to a historic *Crunchbase* snapshot
454 from January 24, 2022 only, about half of the
455 companies in our dataset announced their seed or
456 pre-seed funding round before that date making
457 it theoretically possible that some of the *Crunch-
458 base* information represents a lookahead. However,
459 we consider it unlikely that this information pro-
460 vides substantial information about the future as
461 the potentially affected data fields do not provide
462 any information on future funding rounds. At the
463 same time, our GPT-IO prompting baseline may
464 be affected by look-ahead bias, if GPT-5-mini was
465 trained on data about the companies in our dataset,
466 possibly overstating the performance of this base-
467 line. Long-term studies capturing reliable data and
468 evaluating startup performance over a long time-
469 frame would be needed to certainly rule out any
470 risk of look-ahead bias.

471 As done by several previous studies (see Table
472 4 in the Appendix for an overview), we modeled
473 startup evaluation as a success prediction task with
474 a binary definition of success. Reality is more com-
475 plex, as success may come in various levels and
476 over longer timeframes. Also, the real impact a
477 VC investor would have had on a startup's success
478 after investing in it is in many cases an unobserv-
479 able counterfactual that cannot be assessed in a
480 backtesting study.

481 Lastly, the scope of our analysis is limited to
482 early-stage VC investments (*pre-seed/seed to se-*
483 *ries A*) in European companies and may not gener-
484 alize to other forms of VC or private equity.

485 References

486 Torben Antretter, Ivo Blohm, Dietmar Grichnik, and
487 Joakim Wincent. 2019. *Predicting new venture sur-*
488 *vival: A twitter-based machine learning approach to*
489 *measuring online legitimacy*. *Journal of Business*
490 *Venturing Insights*, 11:e00109.

- Javier Arroyo, Francesco Corea, Guillermo Jimenez-Diaz, and Juan A. Recio-Garcia. 2019. *Assessment of machine learning performance for decision support in venture capital investments*. *IEEE Access*, 7:124233–124243. 491
492
493
494
495
- Shulin Cao, Jiajie Zhang, Jiaxin Shi, Xin Lv, Zijun Yao, Qi Tian, Lei Hou, and Juanzi Li. 2023. *Probabilistic tree-of-thought reasoning for answering knowledge-intensive complex questions*. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 12541–12560, Singapore. Association for Computational Linguistics. 496
497
498
499
500
501
502
- Chi-Min Chan, Weize Chen, Yusheng Su, Jianxuan Yu, Wei Xue, Shanghang Zhang, Jie Fu, and Zhiyuan Liu. 2023. *Chateval: Towards better llm-based evaluators through multi-agent debate*. *Preprint*, arXiv:2308.07201. 503
504
505
506
507
- Harrison Chase. 2022. Langchain. <https://github.com/langchain-ai/langchain>. Accessed: 2025-01-15. 508
509
510
- Kimberly Chong and David Tuckett. 2014. *Constructing conviction through action and narrative: how money managers manage uncertainty and the consequence for financial market functioning*. *Socio-Economic Review*, 13(2):309–330. 511
512
513
514
515
- Francesco Corea, Giorgio Bertinetti, and Enrico Maria Cervellati. 2021. *Hacking the venture industry: An early-stage startups investment framework for data-driven investors*. *Machine Learning with Applications*, 5:100062. 516
517
518
519
520
- Dominik Dellermann, Nikolaus Lipusch, Philipp Ebel, Karl Michael Popp, and Jan Marco Leimeister. 2021. *Finding the unicorn: Predicting early stage startup success through a hybrid intelligence method*. *Preprint*, arXiv:2105.03360. 521
522
523
524
525
- Emily Gavrilenko, Foaad Khosmood, Mahdi Rashtad, and Sadra Amiri Moghaddam. 2023. *Improving startup success with text analysis*. *Preprint*, arXiv:2312.06236. 526
527
528
529
- Paul A. Gompers, Will Gornall, Steven N. Kaplan, and Ilya A. Strebulaev. 2020. *How do venture capitalists make decisions?* *Journal of Financial Economics*, 135(1):169–190. 530
531
532
533
- Shanshan Han, Qifan Zhang, Yuhang Yao, Weizhao Jin, Zhaozhuo Xu, and Chaoyang He. 2024. *Llm multi-agent systems: Challenges and open problems*. *CoRR*, abs/2402.03578. 534
535
536
537
- IMARC Group. 2024. *Venture capital investment market: Global industry trends, share, size, growth, opportunity and forecast 2025–2033*. <https://www.imarcgroup.com/venture-capital-investment-market>. Market size projections 2024–2033. 538
539
540
541
542
543

544	Suprasith Jarupathirun and Fatemeh “Mariam” Zahedi.	599
545	2007. Dialectic decision support systems: System design and empirical evaluation. <i>Decision Support Systems</i> , 43(4):1553–1570. Special Issue Clusters.	600
546		601
547		602
548	Alex Kim, Keonwoo Kim, and Sangwon Yoon. 2024.	603
549	DEBATE: Devil’s advocate-based assessment and text evaluation. In <i>Findings of the Association for Computational Linguistics: ACL 2024</i> , pages 1885–	604
550	1897, Bangkok, Thailand. Association for Computational Linguistics.	605
551		606
552		607
553		608
554	Tian Liang, Zhiwei He, Wenxiang Jiao, Xing Wang,	609
555	Yan Wang, Rui Wang, Yujiu Yang, Shuming Shi, and	
556	Zhaopeng Tu. 2024. Encouraging divergent thinking	
557	in large language models through multi-agent debate.	
558	In <i>Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing</i> , pages	
559	17889–17904, Miami, Florida, USA. Association for	
560	Computational Linguistics.	
561		
562	Ollie Liu, Deqing Fu, Dani Yogatama, and Willie Neiswanger. 2025.	610
563	DeLLMa: Decision making under uncertainty with large language models. In <i>The Thirteenth International Conference on Learning Representations</i> .	611
564		612
565		613
566		
567	Yiming Lu, Yebowen Hu, Hassan Foroosh, Wei Jin,	614
568	and Fei Liu. 2025. STRUX: An LLM for decision-	615
569	making with structured explanations. In <i>Proceedings</i>	616
570	<i>of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers)</i> , pages 131–141, Albuquerque, New	617
571	Mexico. Association for Computational Linguistics.	618
572		619
573		620
574		621
575	Abdurahman Maarouf, Stefan Feuerriegel, and Nicolas Pröllochs. 2025.	622
576	A fused large language model for predicting startup success. <i>European Journal of Operational Research</i> , 322(1):198–214.	
577		
578		
579	OpenAI. 2025. Gpt-5-mini (2025-08-07). Large lan-	623
580	guage model.	624
581	Ekin Ozince and Yiğit İhlamur. 2024. Automating venture capital: Founder assessment using llm-powered segmentation, feature engineering and automated labeling techniques. <i>Preprint</i> , arXiv:2407.04885.	625
582		626
583		627
584		628
585	Joon Sung Park, Joseph O’Brien, Carrie Jun Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. 2023. Generative agents: Interactive simulacra of human behavior. In <i>Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology</i> , UIST ’23, New York, NY, USA.	629
586	Association for Computing Machinery.	630
587		631
588		
589		
590		
591		
592	Jingyuan Qi, Zhiyang Xu, Ying Shen, Minqian Liu,	632
593	Di Jin, Qifan Wang, and Lifu Huang. 2023. The art	633
594	of SOCRATIC QUESTIONING: Recursive thinking	634
595	with large language models. In <i>Proceedings of the</i>	
596	<i>2023 Conference on Empirical Methods in Natural Language Processing</i> , pages 4177–4199, Singapore.	
597	Association for Computational Linguistics.	
598		
599	Andre Retterath. 2020. Human versus computer: Benchmarking venture capitalists and machine learning algorithms for investment screening. SSRN working paper.	635
600		636
601		637
602		638
603	Boris Sharchilev, Michael Roizner, Andrey Rumyantsev,	639
604	Denis Ozornin, Pavel Serdyukov, and Maarten de Rijke. 2018. Web-based startup success prediction. In <i>Proceedings of the 27th ACM International Conference on Information and Knowledge Management</i> , CIKM ’18, page 2283–2291, New York, NY, USA.	640
605	Association for Computing Machinery.	641
606		642
607		643
608		644
609		645
610	U.S. Census Bureau. 2025. Business formation statistics. https://www.census.gov/econ/bfs/index.html . Business applications and formations data, accessed 2025.	646
611		647
612		648
613		649
614	Henning Wachsmuth, Gabriella Lapesa, Elena Cabrio,	650
615	Anne Lauscher, Joonsuk Park, Eva Maria Vecchi,	651
616	Serena Villata, and Timon Ziegenbein. 2024. Argument quality assessment in the age of instruction-following large language models. In <i>Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)</i> , pages 1519–1538, Torino, Italia. ELRA and ICCL.	652
617		653
618		654
619		655
620		656
621		657
622		658
623	Henning Wachsmuth, Nona Naderi, Yufang Hou,	659
624	Yonatan Bilu, Vinodkumar Prabhakaran, Tim Alberdingk Thijm, Graeme Hirst, and Benno Stein. 2017. Computational argumentation quality assessment in natural language. In <i>Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics: Volume 1, Long Papers</i> , pages 176–187, Valencia, Spain. Association for Computational Linguistics.	660
625		661
626		662
627		663
628		664
629		665
630		666
631		667
632	Sichao Xiong and Yiğit İhlamur. 2023. Founder-gpt: Self-play to evaluate the founder-idea fit. <i>Preprint</i> , arXiv:2312.12037.	668
633		669
634		670
635	Sichao Xiong, Yiğit İhlamur, Fuat Alican, and Aaron Ontoyin Yin. 2024. Gptree: Towards explainable decision-making via llm-powered decision trees. <i>Preprint</i> , arXiv:2411.08257.	671
636		672
637		673
638		674
639	Andrew L Zacharakis and Dean A Shepherd. 2001. The nature of information and overconfidence on venture capitalists’ decision making. <i>Journal of Business Venturing</i> , 16(4):311–332.	675
640		676
641		677
642		678
643	Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, Hao Zhang, Joseph E Gonzalez, and Ion Stoica. 2023. Judging llm-as-a-judge with mt-bench and chatbot arena. In <i>Advances in Neural Information Processing Systems</i> , volume 36, pages 46595–46623. Curran Associates, Inc.	679
644		680
645		681
646		682
647		683
648		684
649		685
650		686
651	Kamil Źbikowski and Piotr Antosiuk. 2021. A machine learning, bias-free approach for predicting business success using crunchbase data. <i>Information Processing & Management</i> , 58(4):102555.	687
652		688
653		689
654		690

655 A Dataset Creation

656 For our backtesting experiment, we prepared a
657 dataset of real startups from the watchlists of
658 five VC funds (in the following we will refer to
659 these watchlists in singular). The dataset was cre-
660 ated by extracting and merging data from four
661 sources: (1) the **watchlist** of a real VC, (2) **Crunch-**
662 **base** (crunchbase.com), an online database aggre-
663 gating information about businesses, (3) histori-
664 cal snapshots of the **startup websites** retrieved
665 through the *Internet Archive’s Wayback Machine*
666 (web.archive.org), and (4) **search results**
667 obtained through the Perplexity Sonar API. Table 2
668 includes a summary of all extracted data fields and
669 their origination dates.

670 A.1 Preventing Look-Ahead Bias

671 When working with historic data, it is important to
672 consider which information was and was not avail-
673 able to the VC at the time the investment decision
674 had to be made. Including information about the
675 startups originating from a time after the VC’s deci-
676 sion would constitute a *look-ahead bias* ([Żbikowski](#)
677 and [Antosiuk, 2021](#)). Doing so could put the back-
678 tested method at an unfair advantage compared
679 to the real VC, because it could leak information
680 about the startup’s future success. To prevent such
681 a bias, we carefully filter the information available
682 during backtesting by using historic data snapshots.

683 **Cutoff date** In order for the real VC to partic-
684 ipate in the initial funding round, the investment
685 decision had to be made at some point between
686 the VC becoming aware of the startup (the date the
687 startup was added to the VC’s watchlist) and the
688 announcement of the initial funding round, likely
689 closer to the latter. Therefore, we consider the an-
690 nouncement date of the initial funding round as a
691 cutoff and use it to restrict the information that we
692 include in the dataset:

- 693 • **Watchlist:** We limit the data to startups that
694 were added to the watchlist between January
695 1, 2021 and December 31, 2021 and had not
696 yet raised a *series A* or later by the time they
697 were added.
- 698 • **Crunchbase:** We had access to two *Crunch-*
699 *base* snapshots taken on January 24, 2022 and
700 September 1, 2025. We refer to these as the
701 *historic* and the *current* snapshot, respectively.
702 We use only the historic snapshot to extract

703 predictive features. The current snapshot is
704 used to determine whether a startup turned out
705 successful (i.e., received subsequent funding).

- **Startup websites:** For each startup, we
706 retrieve a historic snapshot of its website
707 through *Wayback Machine* from the latest
708 available date before the announcement of the
709 initial funding round.
- **Search results:** For the *Perplexity Sonar API*,
711 we apply a time filter to return only those
712 search results that originate from a time before
713 the announcement of the initial funding round.

715 A.2 Data Preprocessing

716 To construct the dataset, we started with compa-
717 nies that were added to the VC’s watchlist between
718 January 1, 2021 and January 31, 2021 and systemat-
719 ically enriched this set with data from *Crunchbase*,
720 startup websites, and search results, removing com-
721 panies where such enrichments were not possible.
722 The preprocessing followed four phases: (1) data
723 cleaning, (2) entity matching, (3) label assignment,
724 and (4) enrichment. Table 3 reports the number of
725 companies retained after each step and the corre-
726 sponding share of successful companies.

1. **Watchlist export:** The starting dataset con-
727 tained 3,441 companies from the VC’s watch-
728 list that were between January 1, 2021 and
729 January 31, 2021.
2. **Cleaning:** We cleaned the dataset by remov-
731 ing duplicate entries, *Missed Deals*, and com-
732 panies that were considered for a founding
733 round later than *seed*. We further updated all
734 companies website URLs by sending HTTP
735 requests to the domains listed in the watchlist
736 export and recording the final redirect target
737 as the current domain. Finally, we canonical-
738 ized company names and URLs (lowercasing,
739 removing prefixes such as “www”, unicode
740 normalization) for later matching purposes.
741 After cleaning, the dataset contained 3,357
742 companies.
3. **Entity matching:** To enrich the watchlist
744 records with additional data from *Crunchbase*,
745 we performed a left-join between the cleaned
746 dataset and the current *Crunchbase* snapshot
747 and then the historic *Crunchbase* snapshot.
748 Matching followed a strict precedence: (1)

Source	Data Field	Description	Value As Of
Watchlist	<i>Name</i> F	The company name.	Some time in 2021
	<i>Domain</i> F	The company web domain.	Some time in 2021
	<i>Date Added</i> P	The date the company was added to the watchlist. Only companies added between January 1, 2021 and December 31, 2021 are considered.	Some time in 2021
	<i>Status</i> P	The last stage reached in the investment process (e.g., <i>Added to Watchlist</i> , <i>Initial Review</i> , or <i>Investment Made</i>). Used to determine whether the real VC invested in the company.	Some time in 2021
Crunchbase	<i>Funding Rounds</i> P	A list of all funding rounds, including round type (<i>pre-seed</i> , <i>seed</i> , <i>series A</i> , <i>IPO</i> , etc.), amount, and announcement date of the round.	September 1, 2025
	<i>Current Name</i> P	The company's current name.	September 1, 2025
	<i>Current Domain</i> P	The company's current web domain.	September 1, 2025
Crunchbase	<i>Short Description</i> F	Short description of the company.	January 24, 2022
	<i>Long Description</i> F	Long description of the company.	January 24, 2022
	<i>Industries</i> F	A list of industries the company is operating in.	January 24, 2022
	<i>Team</i> F	The names of the team members, their education, prior work experience, and current roles.	January 24, 2022
	<i>Historic Name</i> P	The company's former name.	January 24, 2022
	<i>Historic Domain</i> P	The company's former web domain.	January 24, 2022
	<i>Historic Funding Rounds</i> P	A list of all historic funding rounds, including round type, amount, and announcement date.	January 24, 2022
	<i>Website</i> F	The historic HTML content of the company's website (only the homepage).	Various dates
Startup Websites	<i>Archived</i> P	The date and time when the website was captured.	Various dates
	<i>Results</i> F	The list of search results for a given query, including title, content snippet, and URL.	Various dates

Table 2: An overview of the data sources and extracted data fields used for the dataset creation. Data fields marked with F are used as predictive features (independent variables), whereas data fields marked with P are only used during preprocessing (e.g., for merging data from different sources) and were **not** made available to the prediction method.

750 exact current domain match, (2) exact historic domain match, and (3) fuzzy name-and-domain match with a similarity threshold of 95%. After matching, the dataset contained
751 1,623 companies.
752
753
754

- 755 4. **Label assignment:** Success labels were
756 constructed from the current *Crunchbase* snapshot.
757 Companies that had raised a series A or later funding round by September 1, 2025
758 were labeled as *successful*, all others as *unsuccessful*.
759
760 5. **Enrichment:** We further enriched each startup’s data with additional historic information from the startup’s website, web search results, and historic *Crunchbase* snapshot to extract predictive features including long and short company descriptions, industry, team setup, website content, and information from online articles. We removed companies where no founding team information was available, leaving 637 companies.
770

771 B DIALECTIC Design Details

772 B.1 Seed Questions

773 To kick off DIALECTIC’s question decomposition,
774 we provide it with a set of seed questions Q_0 . These
775 are intended to guide DIALECTIC’s fact gathering
776 efforts by giving it a rough scaffolding of relevant
777 fact categories. VC investors typically assess
778 startups across the following dimensions: general
779 company, market, product/service, entrepreneurial
780 team, and funding (Retterath, 2020). While rich
781 and reliable funding information is typically pri-
782 vate and was not available to us for every startup in
783 our dataset, we dedicate one seed question to each
784 of the remaining four dimensions. Specifically, we
785 use the following four seed questions:

- 786 1. **General Company:** “How do the company’s
787 sector, development stage, and geography
788 align with the VC’s investment strategy?”
789
790 2. **Team:** “Who are the key members of the
791 founding team, and what relevant experience
792 and track record do they have?”
793
794 3. **Product:** “What are the product’s core fea-
795 tures, underlying technology, and forms of
796 protection?”

- 797 4. **Market:** “What is the current size, historical
798 growth, and forecast growth of the target mar-
799 ket, and which customer needs or market gaps
800 does the company address?”
801
802
803

804 B.2 Generating Question Trees

805 In order to evaluate each startup in more detail,
806 DIALECTIC decomposes each seed question into
807 lower-level questions tailored to the specific in-
808 dustry of the startup. Together, the seed ques-
809 tions and all their lower-level questions are sup-
810 posed to comprehensively cover the information
811 required by the VC investor to make a decision.
812 In order to derive and answer lower-level ques-
813 tions from the seed questions, we adapt the *Prob-
814abilistic Tree-of-Thought Prompting (ProbTree)*
815 approach (Cao et al., 2023). ProbTree uses an LLM
816 to create hierarchical question decomposition trees
817 (HQDTs) and then answer the questions in a post-
818 order traversal using three different answer strate-
819 gies **Open Book** (retrieving information from on-
820 line sources), **Closed Book** (asking an LLM for its
821 internal knowledge), and **Child Aggregation** (de-
822 riving the answer to a higher-level question from
823 the answers to its lower level questions). For each
824 answer strategy, it calculates a confidence score and
825 then probabilistically chooses the most confident
826 answers for each question.

827 For DIALECTIC, we use a simplified adaptation
828 of ProbTree. It first decomposes a given seed ques-
829 tion into a HQDT in a single LLM prompt. Then
830 it performs a post-order traversal through the tree
831 to answer the questions from leaf nodes to the root
832 node. Unlike ProbTree, we use a single answer
833 prompt for each node and therefore forgo the con-
834 fidence estimation. Each prompt includes a company
835 summary (description, tagline, and team details in-
836 cluding education and prior work experience) and
837 optional web data that the LLM can obtain agenti-
838 cally, if it decides to do so, by using a web search
839 tool. We only allow usage of the web search tool
840 for leaf nodes.

841 For the web searches, we provide the LLM with
842 access to the *Perplexity Sonar API*. We limit search
843 results to five, each described by a title and content
844 snippet. As described in Appendix A.1, we restrict
845 search results to those originating from a time be-
846 fore the announcement date of the initial funding
847 round to prevent look-ahead bias.

Preprocessing Stage	Companies	Success Rate (%)
Export from CRM system	3,441	—
Remove duplicates and missed deals	3,404	—
Remove series A investments by VC funds	3,401	—
Match with current funding data	2,192	21.6
Remove companies without cutoff date	1,715	23.4
Remove companies added post-seed announcement	587	18.7
Apply temporal cutoff (before February 28, 2023)	462	22.0
Match with historic Crunchbase snapshot	259	25.1

Table 3: Dataset size and success rate after successive preprocessing steps. The success rate refers to the share labeled *successful* at the corresponding stage.

B.3 Evaluating Arguments

DIALECTIC includes an evaluator agent that assigns a numeric quality score for each argument. We use an LLM judge (Zheng et al., 2023) to evaluate each argument and apply the taxonomy of argument quality proposed by Wachsmuth et al. (2017). Following the instruction design principles by Wachsmuth et al. (2024), we adapt the taxonomy criteria to the VC context. The revised framework explicitly defines the objective of argumentation (informing the investment decision), establishes domain-specific criteria for argument quality, specifies the intended audience (expert VC investors), and incorporates the surrounding decision context (high-stakes financial environments). Our argument quality evaluation scheme includes the following 14 questions:

1. **Local Acceptability:** Are the premises believable and factually plausible given the provided Q&A facts?
2. **Local Relevance:** Do the premises clearly contribute to supporting or rejecting the conclusion about investment?
3. **Local Sufficiency:** Do the premises provide enough support to justify the conclusion?
4. **Cogency:** Does the argument have premises that are acceptable, relevant, and sufficient to support the investment conclusion?
5. **Credibility:** Does the argument make the author appear credible and trustworthy to VC investors?
6. **Emotional Appeal:** Does the argument create emotions that make the VC investors more receptive?

7. **Clarity:** Does the argument use correct and widely unambiguous language as well as avoid deviation from the issue? 877
878
879
8. **Appropriateness:** Is the style of reasoning and language suitable for a professional VC investment discussion? 880
881
882
9. **Arrangement:** Is the argument well-structured, with a logical order of premises and conclusion? 883
884
885
10. **Effectiveness:** Does the argument succeed in persuading the VC investors toward or against investing? 886
887
888
11. **Global Acceptability:** Would most VCs consider it a valid and legitimate argument? 889
890
12. **Global Relevance:** Does the argument meaningfully contribute to resolving the overall investment question? 891
892
893
13. **Global Sufficiency:** Does the argument adequately anticipate and rebut the main counter-arguments from the argument’s stance? 894
895
896
14. **Reasonableness:** Does the argument resolve the issue in a way acceptable to the VC investors, balancing global acceptability, relevance, and sufficiency? 897
898
899
900

The LLM judge scores each argument across the above 14 criteria using a seven-point Likert scale from 1 (low) to 7 (high). To calculate the final argument quality score, we simply sum up all of the 14 individual scores. The judge also produces justifications explaining each score.

907 C Prompts

908 In the following, we report the prompts that we
 909 used for each of our LLM-based agents. These are:

- 910 • **Decomposer** prompt (Listing 1)
- 911 • **Answer Agent** prompt (Listing 2)
- 912 • **Generator** prompt for pro (Listing 3) and for
 913 contra arguments (Listing 4)
- 914 • **Critic** prompt for pro (Listing 5) and for contra
 915 arguments (Listing 6)
- 916 • **Evaluator** prompt (Listing 7)
- 917 • **Refiner** prompt for pro (Listing 8) and for
 918 contra arguments (Listing 9)
- 919 • **Input Output (IO) Prompting** baseline
 920 prompt (Listing 10)

921 Listing 1: Decomposer Prompt

```
922 SYSTEM: You are good at decomposing a complex
923 question into a hierarchical question
924 decomposition tree (HQDT).
925
926 USER: Please generate a hierarchical question
927 decomposition tree (HQDT) with json format
928 for a given question. In this tree, the root
929 node is the original complex question, and
930 each non-root node is a sub-question of its
931 parent.
932
933 Q: How large is the company's market opportunity
934 (TAM, SAM, SOM)?
935 A: {}
936   "How large is the company's market opportunity
937     (TAM, SAM, SOM)": [
938     "What is the Total Addressable Market (TAM)
939     ?",
940     "What is the Serviceable Available Market (
941     SAM)?",
942     "What is the Serviceable Obtainable Market (
943     SOM)?"
944   ],
945   "What is the Total Addressable Market (TAM)": [
946     [
947       "What customer segments are included in the
948       broadest market?",
949       "What is the total number of potential
950       customers?",
951       "What is the total industry revenue across
952       those segments?"
953   ],
954   "What is the Serviceable Available Market (SAM
955     )": [
956     "Which subset of TAM does the company's
957     product or service directly target?",
958     "What portion of customers can realistically
959     be reached given geography, regulations, or
960     product scope?",
961     "What is the annual spending of these
962     customers?"
```

963] ,	
964 "What is the Serviceable Obtainable Market (964
965 SOM)": [965
966 "What portion of SAM can the company	966
967 realistically capture in the next 3-5 years	967
968 ?", <td>968</td>	968
969 "What customer acquisition assumptions	969
970 support this share?",	970
971 "What expected adoption rate drives this	971
972 forecast?",	972
973 "What annual revenue corresponds to this	973
974 achievable market share?"	974
975]	975
976 }	976
977 Q: What is the competitive landscape, and how is	977
978 the company positioned within it?	978
979 A: {}	979
980 "What is the competitive landscape, and how is	980
981 the company positioned within it?": [981
982 "What is the competitive landscape?",	982
983 "How is the company positioned within the	983
984 competitive landscape?"	984
985],	985
986 "What is the competitive landscape?": [986
987 "Who are the direct competitors?",	987
988 "Who are the indirect competitors or	988
989 substitutes?",	989
990 "What are the major trends shaping	990
991 competition in this space?"	991
992],	992
993 "How is the company positioned within the	993
994 competitive landscape?": [994
995 "What is the company's relative pricing	995
996 strategy?",	996
997 "What is the company's market share or	997
998 traction compared to peers?",	998
999 "Does the company occupy a niche or broader	999
1000 category?",	1000
1001 "What barriers to entry protect the company'	1001
1002 s position?"	1002
1003],	1003
1004 }	1004
1005 Q: What is the company's product differentiation	1005
1006 and value proposition?	1006
1007 A: {}	1007
1008 "What is the company's product differentiation	1008
1009 and value proposition?": [1009
1010 "What is the company's product	1010
1011 differentiation?",	1011
1012 "What is the company's value proposition?"	1012
1013],	1013
1014 "What is the company's product differentiation	1014
1015 ?": [1015
1016 "What features or technologies distinguish	1016
1017 the product?",	1017
1018 "How is the product better than alternatives	1018
1019 ?", <td>1019</td>	1019
1020 "What intellectual property (e.g., patents,	1020
1021 proprietary tech) supports defensibility?"	1021
1022],	1022
1023 "What is the company's value proposition?": [1023
1024 "What problem does the product solve for	1024
1025 customers?",	1025
1026 "What measurable benefits (e.g., cost	1026
1027 savings, time savings, revenue uplift) does	1027
1028 it deliver?",	1028
1029 "Why would customers choose this company	1029
1030 over competitors?"	1030
1031]	1031
1032 }	1032

```

1033    ]
1034    }
1035
1036 Here is the question to decompose:
1037 Q: {question}
1038
1039 Generate its HQDT customized for a company in
1040   the {industry} industry.

```

Listing 2: Answer Agent Prompt

```

1042 SYSTEM: Answer the question using company
1043   summary and sub Q&A if provided. Keep answer
1044   concise (<50 words) with data backing.
1045 If unable to answer the question, use web_search
1046   for market data, trends, competitive
1047   analysis, funding info. Focus on industry-
1048   level searches, not specific companies. Use
1049   the tool only if necessary.
1050 Make ONE tool call at a time.
1051
1052 USER: Question: {question}
1053
1054 Company summary: {company_summary}
1055 {qa_pairs}

```

Listing 3: Generator Prompt (Pro Arguments)

```

1058 SYSTEM: You are a very experienced investor at a
1059   top-tier VC fund. You are also a great
1060   storyteller and can tell a compelling story.
1061
1062 USER: Generate {n_pro_arguments} pro arguments
1063   why this company is a good investment
1064   opportunity.
1065
1066 Each argument should be concise (max. 100 words)
1067   and backed by specific data from the
1068   questions and answers.
1069
1070 A good argument provides a unique perspective on
1071   the investment opportunity that addresses
1072   the following criteria:
1073
1074 1. Local Acceptability - Are the premises
1075   believable and factually plausible given the
1076   provided Q&A facts?
1077 2. Local Relevance - Do the premises clearly
1078   contribute to supporting or rejecting the
1079   conclusion about investment?
1080 3. Local Sufficiency - Do the premises provide
1081   enough support to justify the conclusion?
1082 4. Cogency - Does the argument have premises
1083   that are acceptable, relevant, and
1084   sufficient to support the investment
1085   conclusion?
1086 5. Credibility - Does the argument make the
1087   author appear credible and trustworthy to VC
1088   investors?
1089 6. Emotional Appeal - Does the argument create
1090   emotions that make the VC investors more
1091   receptive?
1092 7. Clarity - Does the argument use correct and
1093   widely unambiguous language as well as avoid
1094   deviation from the issue?
1095 8. Appropriateness - Is the style of reasoning
1096   and language suitable for a professional VC
1097   investment discussion?
1098 9. Arrangement - Is the argument well-structured,
1099   with a logical order of premises and
1100   conclusion?

```

10. Effectiveness - Does the argument succeed in persuading the VC investors toward or against investing?
11. Global Acceptability - Would most VCs consider it a valid/legitimate argument?
12. Global Relevance - Does the argument meaningfully contribute to resolving the overall investment question?
13. Global Sufficiency - Does the argument adequately anticipate and rebut the main counterarguments from the argument's stance?
14. Reasonableness - Does the argument resolve the issue in a way acceptable to the VC investors, balancing global acceptability, relevance, and sufficiency?

Here are the questions and answers about the company:
{qa_pairs}

Provide the qa_indices that were used to generate the argument.

Listing 4: Generator Prompt (Contra Arguments)

- ```

1124 SYSTEM: You are a very experienced investor at a
1125 top-tier VC fund. You are also a great
1126 storyteller and can tell a compelling story.
1127
1128 USER: Generate {n_contra_arguments} contra
1129 arguments why this company is a bad
1130 investment opportunity.
1131
1132 Each argument should be concise (2-3 sentences)
1133 and backed by specific data from the
1134 questions and answers.
1135 Lack of data is not a good contra argument.
1136
1137 A good argument provides a unique perspective on
1138 the investment opportunity that addresses
1139 the following criteria:
1140
1141 1. Local Acceptability - Are the premises
1142 believable and factually plausible given the
1143 provided Q&A facts?
1144 2. Local Relevance - Do the premises clearly
1145 contribute to supporting or rejecting the
1146 conclusion about investment?
1147 3. Local Sufficiency - Do the premises provide
1148 enough support to justify the conclusion?
1149 4. Cogency - Does the argument have premises
1150 that are acceptable, relevant, and
1151 sufficient to support the investment
1152 conclusion?
1153 5. Credibility - Does the argument make the
1154 author appear credible and trustworthy to VC
1155 investors?
1156 6. Emotional Appeal - Does the argument create
1157 emotions that make the VC investors more
1158 receptive?
1159 7. Clarity - Does the argument use correct and
1160 widely unambiguous language as well as avoid
1161 deviation from the issue?
1162 8. Appropriateness - Is the style of reasoning
1163 and language suitable for a professional VC
1164 investment discussion?
1165 9. Arrangement - Is the argument well-structured,
1166 with a logical order of premises and
1167 conclusion?
1168 10. Effectiveness - Does the argument succeed in
1169 persuading the VC investors toward or

```

- against investing?
- 11. Global Acceptability - Would most VCs consider it a valid/legitimate argument?
- 12. Global Relevance - Does the argument meaningfully contribute to resolving the overall investment question?
- 13. Global Sufficiency - Does the argument adequately anticipate and rebut the main counterarguments from the argument's stance?
- 14. Reasonableness - Does the argument resolve the issue in a way acceptable to the VC investors, balancing global acceptability, relevance, and sufficiency?

Here are the questions and answers about the company:

{qa\_pairs}

Provide the qa\_indices that were used to generate the argument.

**Listing 5: Critic Prompt (Pro Arguments)**

SYSTEM: You are a very experienced VC investor against investing in the company. However, your colleague thinks it is a good investment opportunity.

Your job is to criticize the pro argument given by your colleague using the questions and answers about the company and defend your position.

Be direct to persuade your colleague not to invest in the company.

USER: Here are the questions and answers about the company:  
{qa\_pairs}

Here is the argument you have to criticize to persuade the colleague not to invest in the company:  
{argument}

Keep your critique concise in 3-4 sentences.

Listing 6: Critic Prompt (Contra Arguments)

SYSTEM: You are a very experienced VC investor in favor of investing in the company. However, your colleague thinks it is a bad investment opportunity.

Your job is to criticize the given contra argument given by your colleague using the questions and answers about the company and defend your position.

Be direct to persuade your colleague to invest in the company.

USER: Here are the questions and answers about the company:  
{qa\_pairs}

Here is the argument you have to criticize to persuade the colleague to invest in the company:  
{argument}

Keep your critique concise in 3-4 sentences.

**Listing 7:** Evaluator Prompt

SYSTEM: You are an impartial LLM judge to evaluate the quality of an argument in the VC investment context. The goal of the argument is to support or reject a startup investment decision in a persuasive way.

The quality of an argument in the venture capital investment context should be evaluated along the following 14 dimensions. For each dimension, assign a score from 1 (Low) to 7 (High), and provide a short feedback (1 sentence) how to improve the score.

## 14 Dimensions:

1. Local Acceptability - Are the premises believable and factually plausible given the provided Q&A facts?
  2. Local Relevance - Do the premises clearly contribute to supporting or rejecting the conclusion about investment?
  3. Local Sufficiency - Do the premises provide enough support to justify the conclusion?
  4. Cogency - Does the argument have premises that are acceptable, relevant, and sufficient to support the investment conclusion?
  5. Credibility - Does the argument make the author appear credible and trustworthy to VC investors?
  6. Emotional Appeal - Does the argument create emotions that make the VC investors more receptive?
  7. Clarity - Does the argument use correct and widely unambiguous language as well as avoid deviation from the issue?
  8. Appropriateness - Is the style of reasoning and language suitable for a professional VC investment discussion?
  9. Arrangement - Is the argument well-structured, with a logical order of premises and conclusion?
  10. Effectiveness - Does the argument succeed in persuading the VC investors toward or against investing?
  11. Global Acceptability - Would most VCs consider it a valid/legitimate argument?
  12. Global Relevance - Does the argument meaningfully contribute to resolving the overall investment question?
  13. Global Sufficiency - Does the argument adequately anticipate and rebut the main counterarguments from the argument's stance?
  14. Reasonableness - Does the argument resolve the issue in a way acceptable to the VC investors, balancing global acceptability, relevance, and sufficiency?

Listing 8: Refiner Prompt (Pro Arguments)

SYSTEM: You are a very experienced investor at a top-tier VC fund. You are sure that the company is a good investment opportunity. Your job is to revise your argument to reach

```

1305 better argument quality scores.
1306
1307 USER: Here are the Q&A facts about the company:
1308 {qa_pairs}
1309
1310 Here is your previous argument:
1311 {argument}
1312
1313 Here are the argument quality scores (1-7) to
1314 your previous argument:
1315 {argument_feedback}
1316
1317 Refine your argument by improving argument
1318 quality scores.

```

**Listing 9: Refiner Prompt (Contra Arguments)**

```

1320 SYSTEM: You are a very experienced investor at a
1321 top-tier VC fund. You are sure that the
1322 company is a bad investment opportunity.
1323 Your job is to revise your argument to reach
1324 better argument quality scores.
1325
1326
1327 USER: Here are the Q&A facts about the company:
1328 {qa_pairs}
1329
1330 Here is your previous argument:
1331 {argument}
1332
1333 Here are the argument quality scores (1-7) to
1334 your previous argument:
1335 {argument_feedback}
1336
1337 Refine your argument by improving argument
1338 quality scores.

```

**Listing 10: Input Output (IO) Prompting Baseline Prompt**

```

1340 SYSTEM: Assuming you are a venture capital
1341 investor, would you invest in the following
1342 company? Respond with only "Yes" or "No".
1343
1344
1345 USER: Questions and Answers for the company:
1346 {qa_pairs}

```

## D Related Work

Table 4 provides an overview of related work, i.e., studies that propose startup success prediction methods based on machine learning. The table also shows the success criteria used by these studies, the accuracy, recall, and precision achieved by the best-performing models, as well as the interpretability approach taken, if any.

| Work                         | Success Criterion                                                                         | Accuracy     | Recall                                  | Precision                           | Interpretability     |
|------------------------------|-------------------------------------------------------------------------------------------|--------------|-----------------------------------------|-------------------------------------|----------------------|
| Arroyo et al. (2019)         | First event in 3 yrs (AC = acquired, FR = funding round, IPO, CL = closed, NE = no event) | Global 82.2% | FR: 40%, AC: 3%, IPO: very low, NE: 95% | FR: 64%, AC: 33%, IPO: 44%, NE: 85% | Feature-based        |
| Žbikowski and Antosiu (2021) | AC, IPO, Series B                                                                         | 85%          | 34%                                     | 57%                                 | Feature-based        |
| Retterath (2020)             | Follow-on round, trade sale, IPO                                                          | 80%          | 80%                                     | —                                   | No mention           |
| Antritter et al. (2019)      | 5-year survival                                                                           | 76%          | 86%                                     | 80%                                 | Feature-based        |
| Sharchilev et al. (2018)     | Series A+ within 1 yr                                                                     | —            | —                                       | 62.6%                               | Feature-based        |
| Gavrilenko et al. (2023)     | Raise Series A+ within 1 yr                                                               | —            | 82.7%                                   | 74.4%                               | Feature-based        |
| Maarouf et al. (2025)        | IPO, AC, or funding                                                                       | 74.3%        | 78.3%                                   | 59.8%                               | Feature-based        |
| Ozinice and Ihlamur (2024)   | IPO/AC/funding >\$500M                                                                    | 66.7%        | 64.7%                                   | 68.8%                               | Persona-based        |
| Xiong and Ihlamur (2023)     | N/A                                                                                       |              | No backtesting                          |                                     | Pro/contra arguments |
| Xiong et al. (2024)          | IPO/AC/funding >\$500M                                                                    | 87.6%        | 27.1%                                   | 37.3%                               | No mention           |
| Corea et al. (2021)          | IPO, AC, or funding                                                                       | —            | —                                       | —                                   | Feature-based        |

Table 4: Comparison of startup success prediction studies: success criteria, predictive performance of the best model, and interpretability.